

# I Want to Predict with Transformer

Reviewing 'Informer: Beyond Efficient Transformer for Long Sequence Time-Series Forecasting'

Sungguk Cha



# Contents

- Introduce *transformer*
- Introduce limitations of prior approaches
- Review ‘*Informer: Beyond Efficient Transformer for Long Sequence Time-Series Forecasting (AAAI21)*’
- Discussion



# Transformer?

Transformer is

- proposed for *seq2seq* task.
- based on *attention* algorithm.




Figure. Attention algorithm

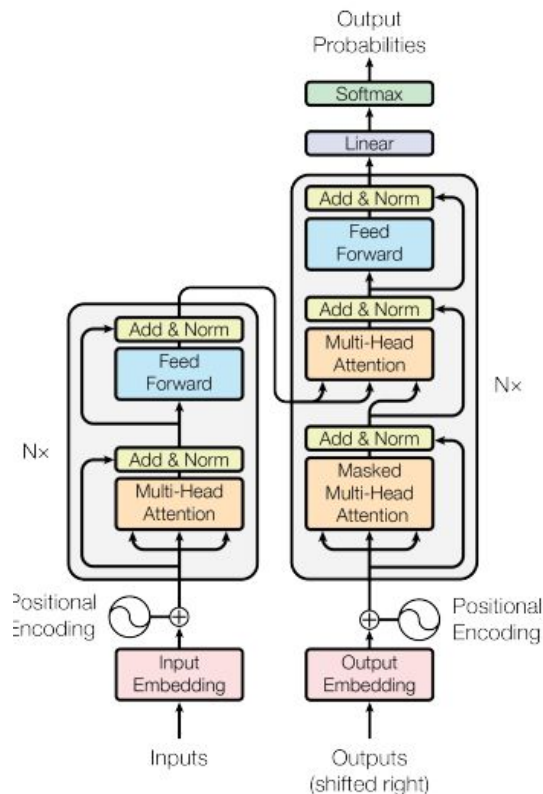


Figure 1: The Transformer - model architecture.

# Transformer Is Overwhelming

Transformer has shown **overwhelming performances** over NLP, CV and multi-modal learning.

NLP: GPT-3 (text generation), BERT (language representation pretraining)

CV: ViT (image recognition)

Vision-NLP: CLIP (zero-shot classification), DALL E (text-to-image synthesis)

Vision-Audio: Lip reading, audio supervised object detection



# Transformer Is Overwhelming

Transformer has shown **overwhelming performances** over NLP, CV and multi-modal learning.

NLP: GPT-3 (text generation), BERT (language representation pretraining)

CV: ViT (image recognition)

Vision-NLP: CLIP (zero-shot classification), DALL E (text-to-image synthesis)

Vision-Audio: Lip reading, audio supervised object detection

In a sense that transformer shows irreplaceable performance on seq2seq,  
**I believe it will be the same over price prediction task.**



# Limitations of Transformers

- The quadratic computation of self-attention:  $O(L^2)$
- The memory bottleneck in stacking layers of  $J$  long inputs:  $O(J * L^2)$
- The speed plunge in predicting long outputs.



# Limitations of Transformers

- The quadratic computation of self-attention:  $O(L^2)$
- The memory bottleneck in stacking layers of  $J$  long inputs:  $O(J * L^2)$
- The speed plunge in predicting long outputs.

I will give you an example of BERT-base model w.r.t. the **space complexity**, while the computation complexity is problematic as well.



# Limitations of Transformers (cont.)

Example with a sequence with 512 tokens.




Figure. Attention algorithm





# Limitations of Transformers (cont.)

Example with a sequence with 512 tokens.

A token is expressed with  $4B * 768 = 3,072B$  (3KB).




Figure. Attention algorithm



# Limitations of Transformers (cont.)

Example with a sequence with 512 tokens.

A token is expressed with  $4B * 768 = 3,072B$  (3KB).

A sequence is  $3,072B * 512 = 1,572,864B$  (1.5MB).




Figure. Attention algorithm



# Limitations of Transformers (cont.)

Example with a sequence with 512 tokens.

A token is expressed with  $4B * 768 = 3,072B$  (3KB).

A sequence is  $3,072B * 512 = 1,572,864B$  (1.5MB).




Figure. Attention algorithm

Attention matrix takes  $512 * 512 * 4B * 768 = 805,306,368B$  (800MB)



# Limitations of Transformers (cont.)

Example with a sequence with 512 tokens.

A token is expressed with  $4B * 768 = 3,072B$  (3KB).

A sequence is  $3,072B * 512 = 1,572,864B$  (1.5MB).




Figure. Attention algorithm

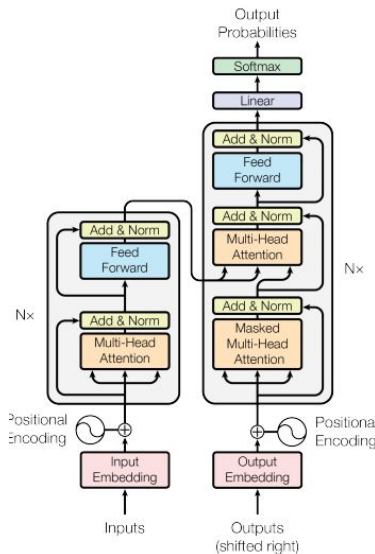


Figure 1: The Transformer - model architecture.

Transformer has 12 layers.  
 $12 * 800 \text{ MB} = 9.6\text{GB}$ .

1 batch requires at least  
about 10GB.

Training is conducted with  
thousands of batches.

I.e., **terabytes GPU** memory  
is needed.

Attention matrix takes  $512 * 512 * 4B * 768 = 805,306,368B$  (800MB)



# Informer!

1. The quadratic computation of self-attention:  $O(L^2)$
2. The memory bottleneck in stacking layers of  $J$  long inputs:  $O(J * L^2)$
3. The speed plunge in predicting long outputs.



# Informer!

## 1. The quadratic computation of self-attention: $O(L^2)$

Heuristic approaches to obtain  $O(L \log L)$  but the efficiency gain is limited.

Sparse Transformer (Child et al. 2019), LogSparse Transformer (Li et al. 2019), Longformer (Beltagy et al. 2020)


Reformer (Kitaev et al. 2019) achieves  $O(L \log L)$  on extremely long sequences.

Linformer (Wang et al. 2020) conditionally obtains  $O(L)$  with risks of degradation to  $O(L^2)$ .



# Efficient Self-Attention

Canonical Self-Attention Takes  $O(L_q L_k)$




$$a' = V_a * Q_a K_a + V_b * Q_a K_b + V_c * Q_a K_c$$

$$\mathcal{A}(\mathbf{q}_i, \mathbf{K}, \mathbf{V}) = \sum_j \frac{k(\mathbf{q}_i, \mathbf{k}_j)}{\sum_l k(\mathbf{q}_i, \mathbf{k}_l)} \mathbf{v}_j = \mathbb{E}_{p(\mathbf{k}_j | \mathbf{q}_i)}[\mathbf{v}_j]$$



# Efficient Self-Attention

Proposed *ProbSparse* Self-attention takes  $O(Lk \log Lq)$



Constrain queries by ignoring sparse correlations.

$$|Q| \rightarrow Lq$$

$$|Q| \rightarrow c \log Lq$$

$$a' = V_a * Q_a K_a + V_b * Q_a K_b + V_c * Q_a K_c$$

$$\mathcal{A}(\mathbf{q}_i, \mathbf{K}, \mathbf{V}) = \sum_j \frac{k(\mathbf{q}_i, \mathbf{k}_j)}{\sum_l k(\mathbf{q}_i, \mathbf{k}_l)} \mathbf{v}_j = \mathbb{E}_{p(\mathbf{k}_j | \mathbf{q}_i)}[\mathbf{v}_j]$$





# Efficient Self-Attention

## Measuring Sparsity

$$p(\mathbf{k}_j | \mathbf{q}_i)$$

If an attention value is closed to  $1 / L$ ,  
it means the attention is **as meaningless as uniform**.



# Efficient Self-Attention

## Measuring Sparsity

$$p(\mathbf{k}_j | \mathbf{q}_i)$$

If an attention value is closed to  $1 / L$ ,  
it means the attention is **as meaningless as uniform**.

$$M(\mathbf{q}_i, \mathbf{K}) = \ln \sum_{j=1}^{L_K} e^{\frac{\mathbf{q}_i \mathbf{k}_j^\top}{\sqrt{d}}} - \frac{1}{L_K} \sum_{j=1}^{L_K} \frac{\mathbf{q}_i \mathbf{k}_j^\top}{\sqrt{d}} \quad \text{measuring KL divergence ....}$$



# Efficient Self-Attention

## Measuring Sparsity

$$p(\mathbf{k}_j | \mathbf{q}_i)$$

If an attention value is closed to  $1 / L$ ,  
it means the attention is **as meaningless as uniform**.

$$M(\mathbf{q}_i, \mathbf{K}) = \ln \sum_{j=1}^{L_K} e^{\frac{\mathbf{q}_i \mathbf{k}_j^\top}{\sqrt{d}}} - \frac{1}{L_K} \sum_{j=1}^{L_K} \frac{\mathbf{q}_i \mathbf{k}_j^\top}{\sqrt{d}} \quad \text{measuring KL divergence ....}$$

Waa	Wba	Wca
Wab	Wbb	Wcb
Wac	Wbc	Wcc

They are to use top-  $c \log L$  correlations only  
according to their importance ranking ( $Q_i K_j$ ).

Computing relatively meaningful ones only.

$$|Q| \rightarrow c \log L$$



# Additional Engineering Contributions

- Reduce memory usage by self-attention distilling.
- Generates sequential output once
  - which can avoid error accumulation.



# Experiment: Baselines

Methods	Training		Testing
	Time	Memory	Steps
Informer	$\mathcal{O}(L \log L)$	$\mathcal{O}(L \log L)$	1
Transformer	$\mathcal{O}(L^2)$	$\mathcal{O}(L^2)$	$L$
LogTrans	$\mathcal{O}(L \log L)$	$\mathcal{O}(L^2)$	$1^*$
Reformer	$\mathcal{O}(L \log L)$	$\mathcal{O}(L \log L)$	$L$
LSTM	$\mathcal{O}(L)$	$\mathcal{O}(L)$	$L$

<sup>1</sup> The LSTnet is hard to present in a closed form.

<sup>2</sup> The  $\star$  denotes applying our proposed decoder.

Table 4:  $L$ -related computation statics of each layer.



# Experiment: Datasets

They experimented on their own dataset.

- Electricity Transformer Temperature (ETT)
- Electricity Consuming Load (ECL)
- Weather



# Experimental Results

Methods		Informer		Informer <sup>†</sup>		LogTrans		Reformer		LSTMa		LSTnet	
Metric		MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
ETTh <sub>1</sub>	24	<b>0.577</b>	<b>0.549</b>	0.620	0.577	0.686	0.604	0.991	0.754	0.650	0.624	1.293	0.901
	48	<b>0.685</b>	<b>0.625</b>	0.692	0.671	0.766	0.757	1.313	0.906	0.702	0.675	1.456	0.960
	168	<b>0.931</b>	<b>0.752</b>	0.947	0.797	1.002	0.846	1.824	1.138	1.212	0.867	1.997	1.214
	336	1.128	0.873	<b>1.094</b>	<b>0.813</b>	1.362	0.952	2.117	1.280	1.424	0.994	2.655	1.369
	720	<b>1.215</b>	<b>0.896</b>	1.241	0.917	1.397	1.291	2.415	1.520	1.960	1.322	2.143	1.380
ETTh <sub>2</sub>	24	<b>0.720</b>	<b>0.665</b>	0.753	0.727	0.828	0.750	1.531	1.613	1.143	0.813	2.742	1.457
	48	<b>1.457</b>	<b>1.001</b>	1.461	1.077	1.806	1.034	1.871	1.735	1.671	1.221	3.567	1.687
	168	3.489	<b>1.515</b>	3.485	1.612	4.070	1.681	4.660	1.846	4.117	1.674	<b>3.242</b>	2.513
	336	2.723	1.340	2.626	<b>1.285</b>	3.875	1.763	4.028	1.688	3.434	1.549	<b>2.544</b>	2.591
	720	<b>3.467</b>	<b>1.473</b>	3.548	1.495	3.913	1.552	5.381	2.015	3.963	1.788	4.625	3.709
ETTm <sub>1</sub>	24	0.323	<b>0.369</b>	<b>0.306</b>	0.371	0.419	0.412	0.724	0.607	0.621	0.629	1.968	1.170
	48	0.494	0.503	<b>0.465</b>	<b>0.470</b>	0.507	0.583	1.098	0.777	1.392	0.939	1.999	1.215
	96	<b>0.678</b>	0.614	0.681	<b>0.612</b>	0.768	0.792	1.433	0.945	1.339	0.913	2.762	1.542
	288	<b>1.056</b>	<b>0.786</b>	1.162	0.879	1.462	1.320	1.820	1.094	1.740	1.124	1.257	2.076
	672	<b>1.192</b>	<b>0.926</b>	1.231	1.103	1.669	1.461	2.187	1.232	2.736	1.555	1.917	2.941
Weather	24	<b>0.335</b>	<b>0.381</b>	0.349	0.397	0.435	0.477	0.655	0.583	0.546	0.570	0.615	0.545
	48	0.395	0.459	<b>0.386</b>	<b>0.433</b>	0.426	0.495	0.729	0.666	0.829	0.677	0.660	0.589
	168	<b>0.608</b>	<b>0.567</b>	0.613	0.582	0.727	0.671	1.318	0.855	1.038	0.835	0.748	0.647
	336	<b>0.702</b>	<b>0.620</b>	0.707	0.634	0.754	0.670	1.930	1.167	1.657	1.059	0.782	0.683
	720	<b>0.831</b>	<b>0.731</b>	0.834	0.741	0.885	0.773	2.726	1.575	1.536	1.109	0.851	0.757
ECL	48	0.344	<b>0.393</b>	<b>0.334</b>	0.399	0.355	0.418	1.404	0.999	0.486	0.572	0.369	0.445
	168	0.368	0.424	<b>0.353</b>	<b>0.420</b>	0.368	0.432	1.515	1.069	0.574	0.602	0.394	0.476
	336	0.381	<b>0.431</b>	0.381	0.439	<b>0.373</b>	0.439	1.601	1.104	0.886	0.795	0.419	0.477
	720	0.406	0.443	<b>0.391</b>	<b>0.438</b>	0.409	0.454	2.009	1.170	1.676	1.095	0.556	0.565
	960	<b>0.460</b>	<b>0.548</b>	0.492	0.550	0.477	0.589	2.141	1.387	1.591	1.128	0.605	0.599
Count		33		14		1		0		0		2	



# Conclusion

- We introduced Transformer and its limitations.
- We simply reviewed some approaches making transformer *lighter*.





# Discussion

- Additionally, transformer is known to require ***extreme-scale*** data.
- Can we make awesome DNN for price prediction except temporarily dominating RL?

